

Final Report for AI based SE

Seah Kim – 20184228

Jeonggwan Lee – 20184400
Nick Heppert – 20186505

Lingjun Liu – 20186473

December 18, 2018

Contents

1	Introduction	2
2	Problem Formulation	3
2.1	What is a Coverage Target?	3
2.2	Solution Representation	3
2.3	Fitness Function	4
3	Methodology	5
3.1	Extending Fitness Calculation	5
3.2	Implementing NSGA-II	5
3.2.1	Non-Dominating Sort with Preference Criterion	6
3.2.2	Sorting by covered targets	6
3.2.3	Combine Operator	7
4	Evaluation	8
4.1	Setup	8
4.2	Results	8
4.3	Failure Analysis	9
4.3.1	Why Does the Capacity Matters?	10
4.3.2	Consequences	11
5	Conclusion	12
A	Github Information	14

Chapter 1

Introduction

The amount of information is exponentially increasing, a robust and delicate index structure also be needed. SQL, which is a language to manage Database Management System(DBMS) is widely used as a tool which arranges effectively indices and relationship between tables, the area of application is very large, such as Electronic Health Record(EHR), customer management, and distribution management. Software developers can get necessary information through query statement in SQL. However, they need to verify their own query statement and check whether unexpected bugs exist. Otherwise, we can't believe the results of a query and beyond lose reliability of SQL. After verifying, we can judge whether this query statement is semantic or not. If the query statement is simple, developers can make test cases manually considering coverage targets. However, if a query statement is complex, this manual making is very inefficient and hard to deal with.

To handle this problem, EvoSQL[1] made test case generation tool covering whole coverage targets for SQL queries using a genetic algorithm(GA), it shows astonishingly covering performance, 98.6%. Their contributions are, first suggesting the definition of test case generation problem for SQL and a “physical query plan”, which is a series of step to solving in a query. Through the physical query plan, they defined a fitness function and implement crossover and mutation as general GA. They presented seeding strategies, which the model has seeding pool to generate new individual, to fit the characteristic of SQL which normally covers string and integer(days). They compared three methods (Random search, Biased Random Search (supported by Seeding strategy), GA), and show how GA can cover coverage targets of various queries.

GA is much better than the other two methods. But we thought any other improvements than GA, and wondered why they didn't use Multi-Objective Optimization(MOO)[2] methods. What we assumed the weakness of GA is that they didn't arrange the order of coverage targets to solve so that we should wait for solving easy coverage target if it is behind difficult coverage targets, and coverage targets don't share their semantic discovery to others because of limitation of single target strategy.

To apply MOO, we expected that our model can be guided by a similar solution from other coverage targets. Therefore, we first applied NSGA-II[3] as a basic MOO method, changed the crowding distance into “sort fronts by covered target” and added combine operator as a minor technique to satisfy the coverage target easily.

Our contribution is an implementation of an MOO approach of test case generation for SQL, analysis & comparison of results and the answer to a suitability of MOO for test case generation for SQL. we successfully implemented the MOO approach, but our conclusion is it is not applicable to test case generation for SQL, because of too much time cost compared to the performance.

We organized contents as follows. Section 2 describes a genetic algorithm in EvoSQL as a baseline, and the representation of GA setting for SQL test case generation. Section 3 presents our representation of MOO setting, our modified model based on NSGA-II. Section 4 we evaluate our model and analyze failures of it. We conclude the paper in section.

Chapter 2

Problem Formulation

2.1 What is a Coverage Target?

To generate test data of SQL queries, the authors of the original paper first decided a test adequacy criterion. Consider the following SQL query:

```
SELECT * FROM "Product" WHERE "Type" = 'Cosmetic'
```

It contains at least two different branches that could be tested:

- (1) When a row contains `Type = 'Cosmetic'`.
- (2) When a row contains `Type != 'Cosmetic'`.

They selected SQLFpc presented by [4]. SQLFpc is a full predicate coverage criterion for SQL queries which considers selection, joining, grouping, aggregate functions, subqueries, case expressions and null values. Given a SQL query, SQLFpc produces coverage targets in SQL formats. A target is covered when a database returns non-empty results after executing it against the generated data.

As an example, SQLFpc would produce two coverage targets for the query above:

- (1) `SELECT * FROM "Product" WHERE ("Type" = 'Cosmetic')`
- (2) `SELECT * FROM "Product" WHERE NOT ("Type" = 'Cosmetic')`

EvoSQL gets coverage targets by communicating with SQLFpc web service. The authors of the original paper already provided all extracted coverage targets from the queries because the extraction is done via the SQLFpc web service. Thus, if those targets are not getting pre-stored, one would have to connect SQLFpc all the time. Therefore, once EvoSQL executed and extracted coverage targets with SQLFpc, it would pre-store those coverage targets and serialize them into one file. Next time you execute EvoSQL, it will not extract targets again instead it would reload your file in your local computer. The extracted targets provided by the authors of original paper are in an old serialization version so we re-serialize all roughly 2000 queries again with SQLFpc web service.

2.2 Solution Representation

In last section, we mentioned that a target is covered when a database returns one single row after executing it. That means the database has the test data we want. In EvoSQL, they defined “Fixture” this object is literally a set of tables, where each of them contains a list of rows. Therefore, a candidate solution is a set of tables $T = \{T1, \dots, Tn\}$, where each table Ti is composed of rows, i.e., $Ti = \{R1, \dots, Rk\}$. Each row contains cells, i.e., $Rj = \{V1, \dots, Vc\}$, where c is the number of columns in Ti .

Fixture is not a database filled with data so we cannot execute target query to test if it covers the target. They transformed fixture into “INSERT” statements and serialized them. They used HSQLDB, a

relational database management system written in Java, first built the table schema and then executed those “INSERT” statements. Therefore, we have a database filled with our test data so now we can execute the coverage target and then got the result whether it covers or not.

2.3 Fitness Function

The authors of the original paper introduced what a query execution plan for a SQL query is. A query execution plan indicates the operations required to process the query and the order by which they need to be performed. As an example in this query, the execution plan consists of two individual steps.

```
SELECT *  
FROM Cars  
JOIN Tires  
ON Cars.tire_id = Tires.id  
WHERE model = 'Ferrari'
```

The first step is that the two tables get joined. The second step is that we filter out some instances. Each step can contain multiple relational algebra operations. In this example, possible step functions are JOINS or WHERE clauses.

Now we have execution plan for SQL queries so we can estimate the fitness of one Fixture. In EvoSQL, they didn't compared two fixtures by a single value. Instead, they measured (1) MaxQueryLevel, which is how many steps can be reached by one fixture (2) step distance, which how far a single fixture is to satisfy the step where the database engine stopped its execution. Therefore, EvoSQL first compared fixtures with MaxQueryLevel. For example, if Fixture A's MaxQueryLevel is bigger than Fixture B, Fixture A is better than Fixture B and vice versa. However, if both MaxQueryLevel is the same, EvoSQL would compare with step distance. If Fixture A's step distance is shorter than Fixture B, Fixture A is better than Fixture B.

Chapter 3

Methodology

It is noticeable from the result of EvoSQL that GA approach outperforms biased search in solving complex queries, whereas biased search excels in simple ones. Although the GA implementation has an initialization step that is similar to what happens in the biased search, the GA spends time calculating the fitnesses and applying the search operators at every iteration of evolution. All these steps do not happen in the biased search. Therefore, we expected that reformulating the single GA approach into using multi-objective strategies could reduce the inefficiency of the original search strategy. By using a multi-objective approach, a population from each iteration would be able to share its semantic discovery and save the usage of unnecessary time budget.

3.1 Extending Fitness Calculation

Each coverage target from a single query is set as an objective to be optimized. For every coverage target, fitnesses of a solution are being calculated at the same time. We newly defined `class FixtureMOO`, which has a member variable of the list containing the fitnesses of targets(`fitness_moo`). We utilized the same fitness function with EvoSQL for effective comparison to single objective approach. Once a `FixtureMOO` object is created, it calculates the fitness of the query on a test.

```
public class FixtureMOO extends Fixture{
    [...]

    private List<FixutreFitness> fitness_moo = new ArrayList<FixtureFitness>();
    [...]

    public int calculate_fitness_moo(List <String> paths_to_test,
    Map <String, TableSchema> tableSchemas){
        [...]
    }
}
```

Listing 3.1: FixtureMOO.java

3.2 Implementing NSGA-II

With multiple objectives in a problem indicates that there is a set of optimal solutions instead of a single optimal solution. Mostly, these solutions are Pareto-optimal and hard to say that one solution to be better




Figure 3.1: Distribution of coverage targets (i.e. objectives) in our whole test query dataset. There were also queries with more than 28 coverage targets but they were truncated as they could be considered as not important.

Table 3.1: Number of Queries by coverage targets

Coverage Targets	0	1-2	3-4	5-6	7-8	9-10	11-15	16-20	21+	Total
	68	645	333	372	308	90	61	26	53	1956

than the other. Among several classical multi-objective evolutionary algorithms, we choose Non-dominated Sorting Genetic Algorithm II (NSGA-II) [3] for our implementation since it is simple and straightforward.

3.2.1 Non-Dominating Sort with Preference Criterion

Over 60 percent of queries that EvoSQL handles have more than three coverage targets (fig. 3.1, table 3.1). It is shown that NSGA-II has been successful in solving the optimization problem with multiple objectives. However, it is not effective for solving problems with *many* objectives. In case of many objectives, most of the individuals are non-dominated by each other and it makes hard setting up a Pareto front.

To overcome this limitation, we applied preference sorting which [2] proposed. [2] suggested a way of considering both the non-dominance relation and the preference criterion. It imposes an order of preference among non-dominated solutions. In short, preference sorting adds individuals which achieved the lowest in each objective to the first non-dominating front. In our case, a set of tables which covers the maximum number of target among other solutions will be located on the first Pareto front.

3.2.2 Sorting by covered targets

As a next step, nondominated solutions are sorted by the number of covered targets. Originally in NSGA-II, comparing the crowdedness among nondominated solutions was introduced. This procedure ensures diversity of solutions. However, the current implementation of EvoSQL does not provide an explicit calculation of the fitness as the paper said. Only the comparison of fitnesses between two individual is possible. Therefore, it is not feasible to get a numeric value for setting up the boundary points. This problem

is still left as an open Github issue.¹

3.2.3 Combine Operator

In last section, we sorted by calculating how many targets one individual covers. We also examined each individual covering which target. We found that one individual covers the target that another individual doesn't cover. Therefore, we came up an idea that if we combine these two individuals, we can get new individual covering all targets that these two individuals cover. The combine operator adds all rows from two fixtures into one fixture. However, the size of new offspring would be two times of its parent. That means this operation would fill up one individual's tables quite fast.

¹<https://github.com/SERG-Delft/evosql/issues/41>

Chapter 4

Evaluation

After successfully implementing our MOOP extension we started to conduct various experiments. In the following we will give an overview of our experimental setup, followed by the results and analyzing them.

4.1 Setup

We used the previous serialized queries (see section 2.1) from the three different projects that were provided with the Github repository² by the authors of the original EvoSQL paper. As previously explained in section 2.1 for each query we extract multiple coverage targets. In fig. 3.1 we report the distribution of all coverage targets from all three projects.

As seen in the data distribution a lot of our queries have two coverage targets which need to be covered in order to say the query is solved. The main reason why so many queries have two targets is because if we have single condition in our query, for instance:

```
SELECT user_id AS 'userId' FROM autofollow WHERE entity_type = 'Account'
```

The two possible targets are:

- (1) SELECT "user_id" AS "userId" FROM "autofollow" WHERE NOT ("entity_type" = 'Account')
- (2) SELECT "user_id" AS "userId" FROM "autofollow" WHERE ("entity_type" = 'Account')

In our experiments we measured the average time it takes to solve a query with a given amount of targets. To ensure we don't run into unsolvable targets we limit the maximum execution time to 30 minutes.

In the case of the original implementation we sum up the amount of time needed to solve each target individually. This sum then forms the total execution time for the single-objective optimization. When measuring the multi-objective optimization we only need to measure the overall execution time.

4.2 Results

To our personal surprise the time to fully cover a query increased for our multi-objective optimization implementation with increasing amount of targets to cover, even hitting time outs when having a lot of coverage targets (objectives).

²<https://github.com/SERG-Delft/evosql>




Figure 4.1: Average execution time for each covered target. For the last data point (15 coverage targets, our MOO was able to only cover 3 targets in the given time frame of 30 mins.)

Table 4.1: Ratio of execution time between original single-objective optimization and our implemented multi-objective optimization (if > 1 : multi-objective optimization is slower)

Coverage Targets	3	4	5	6	7	8	9	10
Ratio (MOO/GA)	2.259	1.587	1.501	1.581	4.265	5.860	6.816	13.463

4.3 Failure Analysis

If we look at table 4.1 we can see that the actual ratio between the two different execution times is not staying constant. This leads to the conclusion on our side that with increasing amount of coverage targets our multi-objective algorithm performs worse and worse.

To further investigate why this happens we took a look at the execution time of each component. Here we could narrow it down that our fitness calculation becomes worse with increasing amount of coverage targets. Initially, this makes sense as we have to calculate the fitness **amount of coverage targets**-times for each individual. But as we have to execute the optimization algorithm only once in the case of multi-objective optimization and not "amount of coverage targets" times it should even out in the end.

Next, we found out that the actual time for calculating the fitness for one coverage target increases as well with having more coverage targets. When looking at this aspect, the only thing that changes with increasing amount of coverage targets is the maximum amount of rows we can have in an individual.

We have to scale the maximum amount of rows based on the amount of coverage targets as our individuals need to have enough "capacity" to potentially cover all targets. We do so by setting the capacity to

```
max_rows = max_rows_per_target * amount_coverage_targets
```

We adapted the given `max_rows_per_target` from the single objective optimization and set it to 4.

4.3.1 Why Does the Capacity Matters?

We will now analyze why the fitness calculation becomes more costly when the amount of rows increase per table.

Theoretical Analysis

Assuming we have a simple example with two tables, if we now execute a query on these two tables with a form like `WHERE table1.field == "Foo" AND table2.field == "Bar"` we have to calculate the minimum fitness for each possible assignments of `table1.field` as well as `table2.field`.

Abstractly this could be described in the following way:

```
min_fitness = Null
foreach row1 in table1:
    foreach row2 in table2:
        fitness = query.get_distance(row1, row2)
        if fitness < min_fitness:
            min_fitness = fitness
return min_fitness
```

So in this hypothetical case, where no optimizations are applied the total amount of query plan analyses (in `get_distance`) can be calculated by

$$c(T) = \prod_{t \in T} |t| \quad (4.1)$$

where T is the set of all tables and the norm $|t|$ is the amount of rows in a table.

For our convenience we set the maximum amount of rows m_r equal for all tables. Therefore we can construct our worst case scenario/an upper bound for eq. (4.1) by assuming all tables are filled completely:

$$c(T) = m_r^{|T|} \quad (4.2)$$

where $|T|$ is the total number of tables.

If we now look at a constructed example with:

- 27 coverage targets,
- 6 tables
- 4 max rows per target

In the single objective case the total amount of comparisons is

$$\underbrace{27}_{\text{Coverage Targets/Single Optimization Executions}} \cdot \underbrace{4^6}_{\text{Equation (4.2)}} \cdot \text{generations} \cdot \text{individuals} \quad (4.3)$$

while in the multi-objective case we construct the amount of comparisons the following way

$$\underbrace{1}_{\text{Optimization Execution}} \cdot \underbrace{(27 \cdot 4)^6}_{\text{Equation (4.2)}} \cdot \text{generations} \cdot \text{individuals}. \quad (4.4)$$

There in this simple example the ratio of comparisons is

$$\frac{\text{Comparisons MOO}}{\text{Comparisons SO}} = \frac{(27 \cdot 4)^6}{27 \cdot 4^6} = 14,348,907 \quad (4.5)$$

Of course this example is constructed and does not reflect all the cases we have in our distribution (see fig. 3.1). But even if we lower the amount of coverage targets to 3, this ratio still would be 243.

Table 4.2: Empirical comparison of the executed query plans

Query	Tables	Targets	Max Rows	Single-Objective		Multi-Objective		Factor/Magnitude	
				Average	Max	Average	Max	Average	Max
Simple	3	11	4	4.82	21.00	1273.99	9246.00	264.21	440.29
Complex	6	27	4	14.17	460.00	16619.06	863391.00	1173.07	1876.94

Empirical Study

In order to check our assumptions we recorded the actual amount of comparisons we did when calculating the fitness. We identified the following piece of code¹ to be crucial for our study:

Listing 4.1: Actual implementation of the previous presented loop

```

for (ComparisonRow c : iterStore.getRows()) {
    try {
        currentDistance = c.getDistance();
    } catch (OperationNotSupportedException e) {
        log.error(e);
        currentDistance = Double.MAX_VALUE;
    }

    [...]
}

```

Here we use the size of `iterStore.getRows()` in order to determine the total amount of executed fitness plans. Compared to the previous nested `foreach`-loops in lis. 4.1 all executed fitness plans are stored in one iterable list `iterStore`.

If we assume our instrumentation doesn't do any optimization and we have full rows the size of the list would be calculated according to eq. (4.1).

We recorded the size of this list for two different queries. The properties and results can be found in table 4.2. Here we can see that luckily our previously calculated ratio of 14,348,907 (See eq. (4.5)) was not fully reached.

4.3.2 Consequences

In our analysis in section 4.3.1 we showed why the maximum row size matters in our specific optimization problem. As we didn't expect that the impact of increasing the maximum row size is that big we couldn't foresee that we would run into such problems.

Nonetheless, to fasten up this issue we tested multiple things. Unfortunately, all without any significant improvement. Please note that the instrumented database application was already operating in "in-memory" memory mode.

First, we tested increasing the cache size of the database, hoping we can have faster access.

Second, we reduced the amount of "back up" writes to the disk.

Sadly, there was no further time to investigate the instrumentation and how to possibly improve it. As the instrumentation was added by the authors of the original paper with small maximum row sizes in mind, it is no surprise that it isn't optimized for relatively large maximum sizes.

¹ See for more context

Chapter 5

Conclusion

The goal of our paper was an implementation of MOO version of EvoSQL, analysis & comparison of results and the answer to a suitability of MOO for test case generation for SQL. In conclusion, we successfully implemented the MOO approach, but it is not applicable to, because the execution of fitness function is too much compared to GA so that the time is out of scope. our model expanded rows to be proportional to the number of objectives, it takes lots of costs. Adding combine operator also increased execution time since the number of rows increased. If we would have more time on this project, we suggested future directions.

First, since our MOO approach with existing fitness function was not suitable, we need to define another numeric fitness function to be used for MOO.

Second, when we find a solution for a particular coverage target, we have kept it and it is a waste of resource. Therefore, if we remove that solved target from the amount of objectives, it reduces the amount of executions.

Third, our basic expectation was, there might be unnecessary time budget to solve infeasible coverage targets if the users didn't eliminate infeasible coverage targets. EvoSQL also removed those infeasible targets manually. We assumed that MOO can handle them without manual elimination since it simultaneously covers multiple coverage targets and detect them if it analyzes the tendency of the non-decreasing fitness function. But, we expected it is hard to judge that a specific coverage target is difficult to solve or infeasible case. Therefore, we need a robust tendency analysis for each coverage target.

Bibliography

- [1] J. Castelein, M. Aniche, M. Soltani, A. Panichella, and A. van Deursen, “Search-based test data generation for sql queries,” in *Proceedings of the 40th International Conference on Software Engineering*. ACM, 2018, pp. 1230–1230.
- [2] A. Panichella, F. M. Kifetew, and P. Tonella, “Reformulating branch coverage as a many-objective optimization problem,” in *Software Testing, Verification and Validation (ICST), 2015 IEEE 8th International Conference on*. IEEE, 2015, pp. 1–10.
- [3] K. Deb, A. Pratap, S. Agarwal, and T. Meyarivan, “A fast and elitist multiobjective genetic algorithm: Nsga-ii,” *IEEE transactions on evolutionary computation*, vol. 6, no. 2, pp. 182–197, 2002.
- [4] C. De La Riva, M. J. Suárez-Cabal, and J. Tuya, “Constraint-based test database generation for sql queries,” in *Proceedings of the 5th Workshop on Automation of Software Test*. ACM, 2010, pp. 67–74.

Appendix A

Github Information

Project Repository: https://github.com/SuperN1ck/aise__project
Usernames:

- jeongggwanlee : Jeongggwan Lee
- Rienshaliu : Lingjun Liu
- bbaa2837 : Seah Kim
- SuperN1ck : Nick Heppert